

service_contracts.md (IronClaw)

This document defines **service boundaries**, **HTTP endpoints**, and **event types/payloads** for the IronClaw microservices architecture.

Principles:

- **Events are append-only** (ledger is the source of truth).
- **Artifacts live in git worktrees** (Vault service owns creation/removal).
- **Workers are stateless** (they write artifacts + AAR, then exit).
- **CO orchestrates** (plans, dispatches, synthesizes).

O) Shared identifiers and conventions

IDs

- garrison_id : string (e.g., "local", "prod-us1")
- theater_id : string (e.g., "demo", "proj_auth")
- run_id : ULID/UUID string
- order_id : ULID/UUID string
- unit_id : string (e.g., "assault_abc123", "fireteam_research")
- worktree_id : string (often same as order_id)
- commit_sha : git SHA string

Status enums

- run_status : OPEN | COMPLETE | FAILED | CANCELLED
- order_status : QUEUED | CLAIMED | RUNNING | BLOCKED | COMPLETED | FAILED | CANCELLED
- severity : LOW | MEDIUM | HIGH | CRITICAL

Time

- Use ISO-8601 UTC: 2026-01-14T16:21:00Z

Response envelopes (recommended)

All services may respond with:

```
{ "ok": true, "data": { ... }, "error": null }
```

On error:

```
{ "ok": false, "data": null, "error": { "code": "...", "message": "..." } }
```

1) Event types (Ledger canonical)

Events are immutable. Consumers derive state from events + snapshots.

Required base event fields

```
{  
  "event_id": "01J...",  
  "ts": "2026-01-14T16:21:00Z",  
  "type": "ORDER_CREATED",  
  "garrison_id": "local",  
  "theater_id": "demo",  
  "run_id": "01J...",  
  "order_id": "01J...",  
  "unit_id": null,  
  "payload": {}  
}
```

Core event types

Run lifecycle

- RUN_CREATED
- RUN_UPDATED (objective/metadata change)
- RUN_COMPLETED
- RUN_FAILED
- RUN_CANCELLED

Order lifecycle

- ORDER_CREATED
- ORDER_ENQUEUED
- ORDER CLAIMED
- ORDER_STARTED
- ORDER_BLOCKED
- ORDER_COMPLETED
- ORDER_FAILED
- ORDER_REISSUED
- ORDER_CANCELLED

Vault/worktree

- WORKTREE_CREATED
- WORKTREE_READY
- WORKTREE_ARCHIVED
- WORKTREE_Removed

Artifacts & reporting

- ARTIFACT_WRITTEN
- AAR_WRITTEN

Integration

- INTEGRATION_READY
- INTEGRATION_STARTED
- INTEGRATION_PASSED
- INTEGRATION_FAILED
- INTEGRATED

Escalation / recovery

- ESCALATION_RAISED
- ESCALATION_ACKED
- RECOVERY_REQUIRED
- RECOVERY_STARTED
- RECOVERY_COMPLETED

System health

- PATROL_TICK (DO/Observer heartbeat)
- SERVICE_DEGRADED
- SERVICE_RECOVERED

2) Artifact formats (shared)

Order file (in worktree)

Path: order.json

```
{  
  "order_id": "order_001",  
  "run_id": "run_123",  
  "theater_id": "demo",  
  "type": "draft",  
  "objective": "Summarize the input into 5 bullets",  
  "inputs": { "text": "..." },  
  "constraints": {  
    "budget_seconds": 60,  
    "max_retries": 1,  
    "tool_policy": { "network": false, "fs_allowlist": ["./"] }  
  },  
  "profile": { "name": "executor_default", "model": "meta-llama/..." }  
}
```

AAR file (in worktree)

Path: aar.json

```
{  
    "order_id": "order_001",  
    "status": "completed",  
    "completed_at": "2026-01-14T16:21:00Z",  
    "summary": "What was done and results",  
    "artifacts": [  
        { "path": "outputs/model_output.txt", "type": "text/plain" }  
    ],  
    "model": {  
        "provider": "io.intelligence",  
        "name": "meta-llama/Llama-3.3-70B-Instruct",  
        "temperature": 0.2  
    },  
    "notes": [],  
    "next_recommendations": []  
}
```

3) Services and contracts

3.1 api_gateway (Chat surface)

Purpose: user-facing entrypoint.

POST /chat Request:

```
{  
    "user_id": "u_123",  
    "theater_hint": "demo",  
    "message": "Build me a summary",  
    "mode": "sync"  
}
```

Response (sync):

```
{  
    "ok": true,  
    "data": {  
        "run_id": "01J...",  
        "reply": "Final user-facing response text",  
        "artifacts": []  
    },  
    "error": null  
}
```

Response (async):

```
{  
    "ok": true,  
    "data": { "run_id": "01J...", "status": "OPEN" },  
    "error": null
```

```
}
```

GET /runs/{run_id} Returns run status + summary.

GET /orders/{order_id} Returns order status + pointers to artifacts/AAR.

3.2 co_service (Commanding Officer)

Purpose: planning, routing, dispatch, synthesis.

POST /runs Create a run from user input (usually called by `api_gateway`). Request: { "user_id": "u_123", "theater_hint": "demo", "objective": "..." }

POST /runs/{run_id}/plan Produces 1-N orders (METE). Response includes created `order_ids` and dependencies.

POST /runs/{run_id}/dispatch Enqueues orders to worker runtime / workflow engine.

POST /runs/{run_id}/synthesize Reads artifacts/AARs and produces final reply.

CO emits events to ledger_service: RUN_CREATED, ORDER_CREATED, ORDER_ENQUEUED, RUN_COMPLETED, etc.

3.3 ledger_service (Events + snapshots)

Purpose: single source of truth for state.

POST /events Append one event. Request:

```
{
  "type": "ORDER_STARTED",
  "garrison_id": "local",
  "theater_id": "demo",
  "run_id": "01J...",
  "order_id": "01J...",
  "unit_id": "assault_abc123",
  "payload": { "attempt": 1 }
}
```

POST /events/batch Append many events at once (recommended for workers).

GET /runs/{run_id} Returns run + derived snapshot.

GET /orders/{order_id} Returns order + derived snapshot + latest known artifact pointers.

GET /health Healthcheck.

Ledger guarantees:

- events are stored durably
- idempotency supported via `event_id` (recommended)

3.4 vault_service (Theater repo + worktrees)

Purpose: create/prepare/cleanup worktrees and write initial operational files.

POST /theaters Create/open a Theater repo. Request: { "theater_id": "demo", "repo_path": "/.../theaters/demo/repo" }

POST /worktrees Create a worktree for an order. Request: { "theater_id": "demo", "order_id": "01J...", "branch": "order_01J..." } Response:

```
{  
  "ok": true,  
  "data": {  
    "worktree_path": "/.../theaters/demo/worktrees/order_01J...",  
    "branch": "order_01J..."  
  },  
  "error": null  
}
```

Emits events: WORKTREE_CREATED , WORKTREE_READY

POST /worktrees/{order_id}/prepare Write order.json , task.md , state.json (optional). Request: { "order": { ... } }

POST /worktrees/{order_id}/archive Archives completed worktree (zip, move, tag) per policy.

DELETE /worktrees/{order_id} Removes worktree safely (git worktree remove).

3.5 worker_service (Assault Units / Fireteams runtime)

Purpose: execute orders, call io.intelligence , write artifacts + AAR. Workers consume jobs from the queue/workflow engine (not necessarily HTTP).

If HTTP-controlled, use: **POST /jobs** Request:

```
{  
  "order_id": "01J...",  
  "theater_id": "demo",  
  "run_id": "01J...",  
  "unit_id": "assault_abc123",  
  "worktree_path": "/.../worktrees/order_01J...",  
  "profile": { "name": "executor_default", "model": "meta-llama/..." },  
  "tool_policy": { "network": false }  
}
```

Worker actions (contract):

1. Load secrets (dotenv / injected)
2. Call io.intelligence (OpenAI-compatible)
3. Write artifacts to outputs/
4. Write aar.json
5. Commit changes
6. Emit events batch: ORDER_STARTED , ARTIFACT_WRITTEN , AAR_WRITTEN , ORDER_COMPLETED (or ORDER_FAILED)

3.6 observer_service (Theater oversight)

Purpose: detect stalls, validate completion, trigger recovery.

POST /theaters/{theater_id}/patrol Runs one patrol tick. Checks:

- Orders stuck in `RUNNING` too long
- Worktrees exist but no recent commits
- Missing AAR/artifacts on “completed”
- Dirty worktree states

Emits:

- `RECOVERY_REQUIRED` (with reason)
- `ESCALATION_RAISED` (severity, route)
- `INTEGRATION_READY` when AAR + artifacts valid

3.7 integration_service (Validation + promotion)

Purpose: apply gates and promote changes from worktree to Theater mainline.

POST /integrate Request:

```
{  
  "theater_id": "demo",  
  "order_id": "01J...",  
  "worktree_path": "/.../worktrees/order_01J...",  
  "gate_profile": "default"  
}
```

Actions:

1. Validate `aar.json` schema
2. Ensure artifacts exist
3. Run configured gates (tests/lint/policy)
4. Promote via merge/cherry-pick

Emits: `INTEGRATION_STARTED` , `INTEGRATION_PASSED` or `INTEGRATION_FAILED` , `INTEGRATED` (with commit SHA)

3.8 do_service (Duty Officer)

Purpose: system-wide patrol and recovery.

POST /patrol One patrol tick:

1. Scan ledger for overdue orders
2. Reissue orders within retry caps
3. Request observer patrols
4. Request worktree quarantine if needed

Emits:

- `PATROL_TICK`
- `RECOVERY_STARTED` , `ORDER_REISSUED` , `RECOVERY_COMPLETED`

- SERVICE_DEGRADED if dependencies failing

3.9 sentinel_service (Watchdog auditor)

Purpose: confirm DO is alive and effective.

POST /audit

- Checks last PATROL_TICK
- Checks backlog trend
- Alerts or triggers failover if needed

Emits:

- ESCALATION_RAISED (CRITICAL) if DO is dead or ineffective

3.10 learning_service (Policy updates)

Purpose: outcome-based routing updates (auditable).

POST /learn/tick Consumes recent events and computes:

- Success rate by profile/task_type
- Latency distribution
- Retry counts
- User correction signals (if logged)

Outputs:

- Updated routing recommendations (policy doc)
- Emits RUN_UPDATED (policy version) or POLICY_UPDATED (optional event type)

4) Queue/Workflow payloads (recommended)

Job message (for worker queue)

```
{
  "job_type": "EXECUTE_ORDER",
  "order_id": "01J...",
  "run_id": "01J...",
  "theater_id": "demo",
  "unit_id": "assault_abc123",
  "worktree_path": "/.../worktrees/order_01J...",
  "attempt": 1,
  "profile": { "name": "executor_default", "model": "meta-llama/..." },
  "constraints": { "budget_seconds": 60, "max_retries": 1 },
  "tool_policy": { "network": false, "fs_allowlist": ["./*"] }
}
```

Job result (emitted as events, not returned)

Use POST /events/batch with ORDER_COMPLETED or ORDER_FAILED plus artifact/AAR pointers.

5) Minimal MVP subset (if you want to cut scope)

To ship MVP with microservices, you can start with only:

1. `api_gateway` (or skip if using CLI)
2. `co_service`
3. `ledger_service`
4. `vault_service`
5. `worker_service`